

Framework Content Generation Rules

The Complete Catalog

Formal, checkable rules — derived from GiGi's own site recommendations — meant to bind any future generator of framework content, human or AI-authored.

- Documentation
- Dashboard
- Work in Progress

Concept stage — not yet part of the framework. Per the Concept → Feasibility Review → Approved lifecycle established in From Concept to Principle, this is Claude's feasibility pass only. Nothing here binds any actual content generation until both Tim and GiGi approve it.

RELATED REQUIREMENTS — REQUIREMENTS REGISTER (GIGI'S EVALUATION DOCS)

- GR-002 — pull key findings out of paragraphs into visual callouts (Designed — Rule 3 + Rule 4)
- GR-004 — modernize typography/spacing (titles, headings, whitespace, width) (Partially Addressed — Rule 5 covers type scale only, not whitespace/width)
- GR-005 — give the tagline more visual prominence site-wide (Partially Addressed — Rule 6 + Rule 12 lock it; only the WIP hub's own masthead is confirmed restyled so far)
- GR-006 — add problem-first navigation entry points (Designed — Rule 9)
- GR-007 — add a simple explanatory sequence diagram (Partially Addressed — Rule 10 locks the sequence's wording; no diagram confirmed built, and Rule 10's "no reordering" is in tension with GR-025)
- GR-009 — replace "optimal" with "designed/appropriate" (Designed — Rule 17)
- GR-010 — replace absolute safety/control claims (Designed — Rule 15)
- GR-011 — reframe the 90/10 HAL 9000 example as origin story (Designed — Rule 18)
- GR-025 — add a "Prepare for it" step to the core design principle (Open — in direct tension with Rule 10's "no reordering permitted"; flagged, not resolved)

This catalog merges what were three separately-numbered documents — Visual System & Hierarchy, Navigation & Structure, and Language & Precision — into one continuously numbered list, currently Rule 1 through Rule 20. Each still turns one of GiGi's site-recommendation groups (or, for Rule 19, a standard Tim set directly) into something checkable, not just a style preference, and each still names what it requires, what it applies to, and, where useful, points at a real page or decision already in the framework that shows the rule in action or in violation. What changes is only the numbering: under three separate "Rule 1...N" sequences, "Rule 3" meant something different depending on which document you were in. Under one catalog, "Rule 14" means exactly one thing everywhere it's cited — in this document, in the compliance audit, or in any future WIP proposal that references it. The three source documents are retained in the WIP folder for history but are superseded by this one.

One related distinction this catalog does *not* collapse: a **rule** is a checkable behavioral constraint on how content gets built ("every artifact draws its colors from the standing set"); a **token** is a specific value a rule governs (the hex code #1B4F8A). Rule 1 below is the rule; the token set it locks is a separate, referenced list, not eighteen more rules of its own. Tokens stay tokens — this catalog just makes sure "Rule N" never needs a second question ("which document's Rule N?") to answer.

Visual System & Hierarchy Rules 1–8

RULE 1 Token Lock

Every framework artifact — HTML page, docx, or PDF — draws its colors and fonts exclusively from the standing token set. Colors: navy #1B4F8A, navy-d #0D2D4F, navy-l #EEF3FA, gray #444, light gray #888, border #DDE3EE, amber #B8873F / amber-bg #FEF3C7 (reserved for draft/pending status only — never used decoratively), and light-blue #9FC2EE / #9FC0E8 (the nav-tagline and role-narrative banner accent). Fonts: Georgia for headlines and serif emphasis, Segoe UI for body and UI text, Consolas for labels, badges, and section numbers — genuine Microsoft Word-compatible fonts only, never Google Fonts or other web-only families, so every artifact renders identically whether it's opened as a live HTML page, a Word document, or a PDF. A generator does not introduce a new color or font for a single artifact; a genuinely new visual need becomes its own proposal, not a local fix.

Applies to: Every new HTML page, docx build, or PDF the framework produces.

Already the de facto practice for colors across the dashboard, docs site, and every docx/PDF built to date — this rule just makes it binding rather than assumed. Corrected during the WIP compliance audit: the original list omitted the light-blue accent, which was already in universal, load-bearing use for nav taglines and role-narrative banners — now folded into the locked set instead of left as an unsanctioned extension. Fonts are the one place this rule is currently ahead of the live framework rather than describing it: every WIP HTML page and docx/PDF now uses this Word-compatible font set, but the live dashboard and docs site still run on the original Google-Fonts-based type system — extending this rule there is a separate, not-yet-made decision.

RULE 2 One Pattern Legend

Anywhere the five HITL patterns (Decision, Knowledge, Document, Transaction, Exception) are referenced beyond a single passing mention, the content uses one fixed color-and-label mapping and either displays or links to the canonical pattern legend. A generator does not invent a new shorthand, abbreviation, or color for a pattern, and does not assume the reader already knows the mapping.

Applies to: Any page or document referencing pattern types — the dashboard, methodology page, maturity/RACI docs, any future role or readiness page.

Real gap today: patterns currently read "more like a legend than building blocks" (GiGi's own words) — no single reusable legend component exists yet.

RULE 3 Extract the Load-Bearing Claim

Before finalizing any section of prose, the generator identifies that section's actual claim — a number, a finding, a governing statement — and renders it separately, in the standing pull-quote/callout treatment, rather than leaving it buried inside a paragraph.

Applies to: Any section built around a specific finding or number.

Reference case: the 18–20 point convergence gap belongs in a callout, the way the HAL 9000 line already does — today it still sits in a sentence.

RULE 4 Callout Budget

No single page or document carries more than three callout or pull-quote boxes. If a draft produces more candidates than that, the generator ranks them by how load-bearing the underlying claim actually is and demotes the rest back to plain prose — it does not keep all of them.

Applies to: Every generated page or document, without exception.

A self-check worth running on our own recent work: an early draft of Talking Through Your Questions had four boxed callouts on one page before this rule existed.

RULE 5 Standing Type Scale as Default

New content starts from the current, already-downsized type scale as its baseline — masthead H1 22–28px, H2 16–20px, body 13–14.5px, table header 11px, table body 12px, caption 12px — never a full-size print or Word default that then needs its own downsizing pass later.

Applies to: Any new document type: the Readiness Table, Trust-dimension work, remaining domain RACI docs, People-side extension pages.

The August 5 downsizing pass across all 13 maturity docs and 7 RACI docs was this rule, applied reactively after the fact — the point is to stop needing that pass at all.

RULE 6 Protect the Signature Line

Wherever the framework masthead appears, "Where human judgment belongs" is set on its own line, visually distinct in weight or color from the framework name beside it — never merged into the same sentence or run together with other boilerplate text.

Applies to: Every masthead: the live dashboard, the docs site, every WIP page.

RULE 7 The Executive Layer Never Replaces the Technical Layer

Any simplified, executive-facing version of technical content — a briefing, a summary card, a pull-statement layer — links back to its full technical source. A generator may summarize; it may never publish a summary as a substitute with no path back to the underlying rigor.

Applies to: The CHRO-facing briefing series, and any future executive-summary layer built on top of existing technical pages.

This is the concrete, already-designed answer to GiGi's own Question 12 — not yet shown to her.

RULE 8 Bottom-Align Column Headings

In any table, a column heading's text vertically aligns to the bottom of the header row, never the middle or top — so a short one-line label and a longer label that wraps to two lines still share the same baseline, and the header row reads as one clean line into the data below regardless of which columns wrap.

Applies to: Every table the framework builds — rule-gap and per-rule findings tables, per-document snapshots, composition/RACI/maturity tables, the WIP catalog.

Corrected during the WIP compliance audit: browsers and Word both default a table header cell to middle/top alignment, not bottom, so this was a real gap rather than an assumed default — fixed across every WIP HTML and docx table with a real column-heading row (the CHRO/HRBP Manager/Recruiter composition tables, From Concept to Principle's and The Shock Wave Effect's tables, and this audit's own tables) as this rule was written.

Navigation & Structure Rules 9–11

RULE 9 Problem-First Entry Points, Additive

Primary navigation offers at least one problem-framed entry point — "I am a CHRO — show me HR," "I am a CFO — show me Finance," "I am redesigning a role — show me Role Transition," "I want to understand AI readiness — show me Maturity" — alongside, never instead of, the existing structural/taxonomy navigation. A generator building a new landing surface includes a problem-first path into the content; it does not remove or bury the structural path to make room for it.

Applies to: The docs site's primary landing/navigation surface; any new top-level page.

Right now the site invites a visitor to "study the architecture" first — the fix adds a second door in, it doesn't replace the first one.

RULE 10 One Canonical Explanatory Sequence

The five-step sequence — Technology Capability → Work Design → Human Judgment → Accountability → Readiness — is the framework's single canonical "explain it in a few seconds" statement. It appears verbatim, in this order, wherever the site attempts to summarize the whole framework in one line. A generator does not paraphrase, reorder, or invent a competing summary sequence.

Applies to: The site's homepage/landing hero, the About page, any executive briefing or summary card.

The same line already got the boxed step-flow treatment in the Talking Through Your Questions document — this rule makes that the one and only version, not a first draft of several.

RULE 11 Hierarchy Fix, Not Redesign

Navigation and structure changes are evaluated first as hierarchy and information-architecture adjustments — reordering, re-labeling, re-emphasizing what already exists — tested against whether a first-time CHRO or CFO reader would know what to click. A full visual redesign, or a new page type built from a blank layout, is out of scope for a navigation fix and requires its own separate proposal.

Applies to: Any navigation or information-architecture change to the docs site or dashboard.

GiGi's own framing, kept as the rule's test: "the foundation is strong; the fix is hierarchy, not a redesign."

RULE 12 Protect the Proven Lines

"Where human judgment belongs" and "HITL as an organizational design standard, not a safety checkpoint" are locked language. A generator does not rewrite, replace, or paraphrase either line — under any editing pass, for any reason — without a separate, explicit change proposal reviewed on its own.

Applies to: Any page carrying these lines today, and any new page restating the framework's core identity.

RULE 13 Name the Organizational Framing Explicitly

When introducing or summarizing the framework, content states plainly that this is an organizational design problem, not only a technology problem — language equivalent to "This is not only a technology problem. It is an organizational design problem." A generator does not let a technology-only framing stand uncorrected in introductory copy.

Applies to: Any About/overview/introductory section, and the opening of any new briefing.

RULE 14 Modeled vs. Observed, Always Labeled

Any statistic or finding presented as evidence — the convergence gap, per-pattern deltas, any future readiness or trust metric — is explicitly labeled modeled/estimated or observed/validated at the point it's stated. A generator never lets a modeled figure read as an empirical fact by omission.

Applies to: Every numeric claim, in every doc, dashboard, and briefing.

The 18–20 point convergence gap is modeled, not observed — already true internally, it just needs that label attached everywhere the number appears, not only where it happens to be caveated today.

RULE 15 No Absolute Claims About Safety or Control

A generator does not use absolute language — "structurally impossible," "nothing is a black box," or equivalents — about risk, safety, or control. It uses calibrated language instead: failure made harder to occur, easier to detect, easier to challenge and contain; decision ownership and escalation staying visible. Calibrated language is honest about residual risk; absolute language overclaims its absence.

Applies to: Any risk-, safety-, or governance-adjacent content.

RULE 16 H% Carries Its Disclaimer

Any time H% or a human-involvement percentage is introduced or displayed, the content states — or links to a statement, reachable from the same view — that H% is not a headcount-reduction metric and should not be read as "X% fewer people." A generator does not present an H% figure in isolation without that context nearby.

Applies to: The dashboard, any RACI/maturity doc showing H%, any executive summary citing the convergence gap.

RULE 17 "Designed," Not "Optimal," Unless Earned

A generator defaults to "designed" or "appropriate" rather than "optimal" when describing a maturity target or reference curve. "Optimal" is used only when a specific, justified claim is actually being made that the target is a universal standard — and that claim is stated outright, not implied by word choice alone.

Applies to: Any reference to M3 / Design Intent, or any future maturity target.

RULE 18 Frame HAL as Origin Story, Not Operating Target

The HAL 9000 90/10 example is framed as describing how the framework was co-created — its origin story — never as a universal operating target for an adopting organization. Wherever the example appears, it carries language equivalent to "HAL was the warning. This framework is an attempt to build the answer," rather than presenting the 90/10 split as a prescriptive ratio.

Applies to: Any page or briefing using the HAL example.

WIP Proposal Structure Rule 19

RULE 19 One Shape for Every WIP Proposal

Every WIP catalog item that proposes a new concept, standard, or change to the framework follows the same five-section structure, in this order: **Why This Now** (what prompted this — a real trigger, not a manufactured one), **The Problem** (what's wrong with the current state, with a concrete example or mock-up where the proposal is visual), **The Recommendation** (the proposed fix, shown the same concrete way as the problem), **Proposed Tokens** (any new values the proposal introduces — colors, names, thresholds — outside the current locked set, noting whether they touch Rule 1 Token Lock or Rule 2 One Pattern Legend), and **Build Note** (what implementing this would actually require, once approved). A generator does not invent a different section order, rename a section, or add a sixth; a section with nothing real to say is dropped entirely rather than padded out.

Applies to: Any new WIP catalog item proposing a concept, standard, or change. Reference-only WIP items that propose nothing — GiGi's own review notes, the compliance audit — aren't proposals and don't need this shape.

Already the shape used by Work Pattern Color Standard, built the day before this rule existed; this rule makes that the standard for every proposal going forward rather than something incidental to that one. H% Provenance Badge predates this rule too and uses its own section names (Why This Now / The Real Distinction / Placement Decided / Naming Standard / Build Note) — a real, not-yet-closed gap, left as-is unless Tim wants it reshaped to match.

Working-Relationship Classification & Progression Rule 20

RULE 20 The Progression Ceiling Rule

When a working-relationship classification's Future State and Progression Ceiling are reasoned at the individual-walkthrough level of rigor (per *Working_Relationship_Progression_Methodology.docx*), that reasoning stands as-is. Where volume makes an individual walkthrough impractical, a generator applies this disclosed heuristic instead, consistently: (1) Current State = Automation is already the floor — Future = Ceiling = Automation. (2) Current State = Judgment set by a standing accountability duty (strategic authority, EEOC/adverse-impact liability, ERISA fiduciary duty, NLRA/union exposure, confidentiality/duty-of-care) is already the ceiling — Future = Ceiling = Judgment. (3) For

Agent-delegation, Augmentation, or Oversight: if a duty attaches specifically to that tier, that tier is the ceiling; absent a flagged duty, the tier reflects a maturity gap and the ceiling moves one notch more automated (Oversight → Agent-delegation, Augmentation → Agent-delegation, Agent-delegation → Automation). A generator does not invent a Future State or Progression Ceiling value by any method other than these two, and never leaves the reasoning implicit.

Applies to: Any classified PCF activity's Future State / Progression Ceiling fields, in the sandbox today or in any future production working-relationship generator.

Applied Sept 16, 2026 to 73 previously-unreasoned PCF 7.0 activities plus the 14 newly-classified 7.5 items — the rule's first real stress-test against live classification data, in the PCF 7.0 sandbox. Documented in full, alongside the five individually-walked-through worked examples it's a lighter-weight substitute for, as Section 11 of Working-Relationship Progression — A Design Methodology.

Any future rule joins this catalog at Rule 21 and onward, whichever category it belongs to — the numbering never restarts, so a citation made today stays correct as the catalog grows.

Tim King & GiGi Garcia | Concept — For Joint Review | September 2026

[Download PDF ↓](#)

[← Back to WIP hub](#)

Timothy P. King · Human Enterprise Architect

Private review area — not part of the published framework

hitldrivenarchitecture.com