

Human Operating Readiness Instrument v0.1

Working methodology + first HRBP Manager / employee grievance pressure test

My goal with this draft is simple: before we let an activity move further into an AI-supported working relationship, I want us to be able to say what has to be true of the people and the organization that will actually live with that change.

This is a working draft for pressure testing. It is not a scorecard, not an employee sentiment survey, and not a final framework standard.

Based on the current Working-Relationship Progression methodology, Requirements Register decisions, and the HRBP Manager / Recruiter / CHRO pilot findings.

1. Why I Think GR-012 Matters

The framework is now much clearer about three separate questions: where an activity is today, how far it should be allowed to progress, and what evidence supports the technical transition. GR-012 adds the piece I do not want us to skip: are the people and the organization actually ready to operate the new relationship once the design changes?

For me, readiness is not whether people are excited about AI. It is whether the human side of the operating model is real. Does the person know what remains theirs? Can they challenge the system? Do they have the authority, skill, time, evidence, and support to intervene when needed? Is the manager prepared? Is the organization asking them to carry a responsibility that looks good on paper but is impossible in practice?

If we can answer those questions consistently, the readiness layer becomes more than an HR overlay. It becomes a real gate in the methodology.

2. Core Design Principle

The question I want this instrument to answer is:

Can the organization operate this proposed human-AI relationship with real authority, capability, evidence, capacity, and escalation?

A transition should not be considered ready simply because the technology performs well. A person can technically remain “in the loop” and still have very little meaningful control if they do not have enough information, authority, skill, time, psychological safety, or organizational support to act.

That is the distinction I want GR-012 to protect: human presence is not the same as human control.

3. What This Instrument Is - and Is Not

It is	It is not
A readiness check for one specific proposed transition.	A generic “AI readiness” score for the whole company.
An evidence-based view of the human operating conditions around the transition.	A sentiment survey about whether people like or trust AI.
A way to surface gaps, controls, and required interventions before progression.	A substitute for technical performance testing.
A repeatable method that should work across roles.	A custom questionnaire that gets reinvented for every job.

4. Unit of Assessment

GR-012 should always be applied to a specific change in the working relationship. I do not want us asking, “Is HR ready for AI?” That is too broad to be useful.

The assessment should read more like this:

- Activity: Manage employee grievances.
- Current relationship: human judgment with AI assistance limited or absent.
- Proposed relationship: AI handles intake, precedent research, case preparation, and documentation; the HRBP Manager retains interpretation, determination, escalation, and accountability.
- Readiness question: what must be true of the HRBP Manager and surrounding organization for that relationship to work in practice?

5. Human Operating Readiness Dimensions

I would start with ten dimensions. I do not see these as ten equal scores. I see them as ten places where a transition can either hold up in the real world or quietly fail.

1. Role clarity

Does the person understand what AI is allowed to do, what stays human, what must be reviewed, what must be decided, and what cannot be delegated?

What I would look for: Role description, workflow, decision-rights language, RACI, training, scenario walkthrough.

What would concern me: The person cannot explain who owns the outcome or assumes AI owns a decision that is still human.

2. Decision authority

When the person believes the AI-supported result is wrong, what are they actually allowed to do?

What I would look for: Override rights, system permissions, escalation process, approval matrix, governed decision points.

What would concern me: The person is expected to review but cannot stop, reject, change, or escalate the outcome.

3. Capability / skill

Does the person have enough domain expertise and judgment to evaluate the AI-supported work independently?

What I would look for: Experience, case review, scenario testing, manager observation, qualifications, exception-handling history.

What would concern me: The reviewer cannot reliably recognize a bad, incomplete, or inappropriate output.

4. Confidence / decision fluency

If the person and the system disagree, will the person make and defend an independent decision?

What I would look for: Scenario exercises, interviews, observed override behavior, manager feedback.

What would concern me: The person has authority but consistently defers to the system because it appears more certain or more data-driven.

5. Psychological safety

Can the person say “the system may be wrong” and act on that concern without unreasonable personal or professional risk?

What I would look for: Speak-up history, employee interviews, leadership behavior, escalation norms, incident reviews.

What would concern me: A formal override exists, but people believe using it will hurt their standing or performance.

6. Capacity / workload

Does the new relationship create enough real capacity for meaningful review, exception handling, and escalation?

What I would look for: Caseload, review time, exception volume, backlog, staffing, verification effort, context switching.

What would concern me: Oversight exists on paper, but workload makes substantive review unrealistic.

7. Manager readiness

Does the manager know how to lead this new operating model rather than supervise the old job with AI added to it?

What I would look for: Manager training, monitoring process, coaching expectations, escalation criteria, performance measures.

What would concern me: The role changes but the manager has no way to coach, monitor, or recognize when the transition is failing.

8. Leadership support / governance

Who owns the outcome, and what happens when the new human-AI relationship fails?

What I would look for: Decision rights, risk tolerance, stop conditions, escalation routes, named owners, governance standards.

What would concern me: Accountability is unclear or no one has authority to pause or reverse the operating model.

9. Trust / evidence quality

What evidence justifies asking the person to rely on the system at this level?

What I would look for: Audit results, error rates, exception data, validation history, known limitations, real-volume testing.

What would concern me: People are expected to rely on a system whose performance has not been demonstrated for the proposed transition.

10. Change capacity

Even if this is the right transition, is this the right time for the organization to make it?

What I would look for: Competing initiatives, staffing changes, implementation calendar, leadership capacity, change fatigue.

What would concern me: The design is sound but the organization cannot absorb another major operating change right now.

6. How I Would Rate Readiness in v0.1

I would avoid a single readiness percentage at this stage. A strong average can hide one serious failure. If a person has no real authority to override an AI-supported decision, nine other good scores should not wash that out.

Finding	Meaning
GREEN - Ready	The condition is in place and supported by evidence.
YELLOW - Conditional	The transition may proceed only if a named gap or control is resolved.
RED - Not Ready	A material gating condition prevents responsible progression right now.
GRAY - Insufficient Evidence	We do not know enough to make a defensible finding.

7. Gating Logic I Want to Pressure-Test

I would treat this as a working hypothesis, not a locked rule. Some dimensions feel close to hard gates because if they fail, the human control we think we have may not actually exist.

Potential hard gates	Usually remediable / conditional
Decision authority Capability / skill Psychological safety Capacity / workload Leadership / governance Trust / evidence quality	Role clarity Confidence / decision fluency Manager readiness Change capacity

I expect the pilots to challenge this split. For example, manager readiness may turn out to be a hard gate in some transitions. That is exactly why I want to test the instrument manually before we build scoring or automate it.

8. First Pressure Test: HRBP Manager / Employee Grievances

I chose employee grievances because the existing pilot already gives us enough structure to test the instrument without pretending we have organization-specific evidence that we do not actually have.

The working scenario is: AI takes on the administrative shell - intake, precedent research, case preparation, and documentation. The HRBP Manager keeps the determination, escalation judgment, and accountability.

The existing methodology already suggests three things worth testing: the remaining human work may become more intense, the HRBP Manager's ability to override AI-supported work is not yet measured, and the practical escalation path points to the Director of HR Ops.

Dimension	Initial working read	What we know / what the model suggests	What I still need before a real verdict
Role clarity	YELLOW	The design distinction is conceptually clear: AI supports the administrative work; the HRBP Manager retains the judgment.	Org-specific workflow and role language showing the boundary in practice.
Decision authority	GRAY	The HRBP Manager is the accountable human in the reference model, but actual override power is not measured.	Proof that the role can reject, pause, override, and escalate AI-supported recommendations.
Capability / skill	GRAY	The role is judgment-heavy, but accountability does not prove the person has the skill to evaluate AI-prepared material.	Case-based evidence or scenario testing showing the person can recognize a bad precedent, weak reasoning, or missing context.
Confidence / fluency	GRAY	No evidence yet on whether the person will challenge a confident-looking AI output.	Scenario exercise or observed behavior under disagreement.
Psychological safety	GRAY	No evidence yet on whether slowing or challenging an AI-supported process is actually supported.	Employee and manager evidence that challenge and escalation are safe in practice.
Capacity / workload	YELLOW	The design suggests a verification tax: less administrative production may create more review, exception, and escalation work.	Caseload, review time, exception rate, and whether freed capacity is real or simply reallocated.
Manager readiness	GRAY	The escalation path identifies Director of HR Ops, but that does not tell us whether the manager is prepared to coach and intervene in the new model.	Manager expectations, escalation criteria, monitoring, and coaching readiness.
Leadership / governance	YELLOW	There is a governance concept and a named accountable role, but the organization-specific operating rules are still unknown.	Named stop conditions, escalation rights, and clarity on who can pause or reverse the transition.
Trust / evidence quality	GRAY	No real-volume evidence has been supplied for AI-supported precedent research or grievance preparation.	Validation results, known failure modes, error data, and examples at real operating volume.
Change capacity	GRAY	No organization-specific information is available.	Competing changes, staffing pressure, implementation timing, and leadership capacity.

9. My Initial Pilot Read

I would not call this organization Ready or Not Ready from the information we have. We do not have the organization-specific evidence needed to make that claim.

My working conclusion is: the proposed relationship is reasonable enough to continue pressure-testing, but the real readiness verdict is currently Insufficient Evidence.

That is not a weakness in the instrument. It is actually the result I want the instrument to be able to produce. It should be comfortable saying "we do not know yet" instead of turning assumptions into a readiness score.

The most important gaps this pilot surfaces are:

- Can the HRBP Manager actually override or stop an AI-supported recommendation?

- Does the role have enough time to perform real review once verification and exception work are counted?
- What evidence is good enough to rely on AI-prepared grievance material?
- Can the person challenge the system without pressure to keep the process moving?
- Is the Director of HR Ops ready to coach and absorb escalations under the new model?

10. Proposed Assessment Record

Every GR-012 assessment should leave behind the same short record:

Field	What gets recorded
Activity / transition	Specific activity, current relationship, proposed next relationship, progression ceiling.
Roles affected	Responsible role(s), accountable role, manager / escalation path where relevant.
Dimension finding	Green / Yellow / Red / Gray.
Evidence	What we actually know and where it came from.
Gap	What is missing or uncertain.
Required action / control	What must happen before progression.
Owner	Who is responsible for closing the gap.
Version / timing	Assessment date, evidence date, reference-model version, review trigger.

11. Questions I Want to Put in Front of Tim

- Do we agree that GR-012 is assessing human and organizational operating conditions, not re-testing the technology?
- Which of the ten dimensions are true gates, and which are diagnostics that can remain Yellow with a mitigation plan?
- Should a High/Critical consequence activity automatically use a weakest-link rule?
- Where should Human Agency sit: inside GR-012, inside the Human Impact Pathway, or as a shared construct used by both?
- What evidence do we want the sandbox to store for each readiness finding? I do not want only a color or score.
- How should we handle multiple Responsible roles - weakest link, separate findings, or role-specific readiness views?
- What would count as enough evidence to say GR-012 'held without modification' across HRBP Manager, Recruiter, and CHRO?

12. What I Do Not Want to Build Yet

- A single 0-100 readiness score.
- A giant employee survey.
- A new readiness instrument for every role.
- A dashboard before the questions and evidence logic have survived manual pressure testing.
- A rule that converts Gray / unknown into Green, Human-only, or any other inferred finding.
- Another construct unless it changes a real progression, governance, role, or intervention decision.

13. Next Working Session

I would use the next session to work one transition all the way through instead of reviewing a long list of requirements.

Step	Working-session goal
1	Confirm the boundary between Technical/Process readiness and GR-012.
2	Run all 10 dimensions against HRBP Manager / grievance work.
3	Decide which findings are hard gates versus conditional.

Step	Working-session goal
4	Agree what evidence gets recorded for each finding.
5	Define where Meaningful Human Control / Human Agency sits.
6	Decide what we need to test next on Recruiter and CHRO.

14. Note I Would Send Tim

Tim - I started turning GR-012 into an actual Human Operating Readiness instrument rather than another concept. I am using ten dimensions, but I am deliberately not averaging them into a score yet. Each dimension produces a finding, the evidence behind it, the gap or control that still needs to be addressed, and the readiness implication. The overall outcome is Ready, Conditional, Not Ready, or Insufficient Evidence.

I also ran a first pressure test against HRBP Manager / employee grievances. The biggest value for me is that the instrument is already surfacing things the technical progression logic does not answer on its own - override authority, verification capacity, psychological safety, manager readiness, and the evidence needed before we ask someone to rely on AI-prepared work.

My suggestion for our next session is that we work this one transition all the way through and decide which dimensions are true gates, what evidence we actually want to capture, and where Human Agency should sit in the architecture before we put anything into the sandbox.

WORKING DRAFT - PRESSURE TEST BEFORE ADOPTION