

EXTERNAL REVIEW · GIGI GARCIA TO TIMOTHY KING · SEPTEMBER 2026

Reflections on Role Maturity, Human Consequence, and the Next Design Layer

Working discussion document — capturing the questions and design implications that surfaced after reviewing the HITL maturity assessment, the role-level dashboard, and the “shock wave” observation about how change moves from the role through the person, manager, team, and accountability.

Documentation

Dashboard

Work in Progress

GiGi's own working reflections, sent to Tim September 2026 ahead of a working session, transcribed here in the site's visual language so they can be read and referenced alongside the rest of the framework. The words and observations below are hers, not Tim's or Claude's — this page is a faithful transcription of her document, not a response to it.

↓ Download PDF

Purpose: capture the questions and design implications that surfaced for me after reviewing the HITL maturity assessment, the role-level dashboard, and your “shock wave” observation about how change moves from the role through the person, manager, team, and accountability. This is not intended as a critique of the current framework. It is an attempt to identify what the next layer may need to explain if we want the model to be both technically rigorous and executive-useful.

1 What Clicked for Me

The more I sit with the HITL maturity work, the more I think the framework is doing two different things — and that the distinction matters.

First, it gives us a structured view of where human involvement remains across processes and roles as AI maturity increases. Second, it begins to reveal where human judgment becomes more concentrated, more consequential, or harder to absorb even as the overall amount of human participation declines.

The framework currently measures human presence. I think the next layer may need to measure human consequence.

That is the central idea I keep coming back to.

Across the HR domain, the current assessment shows aggregate human involvement moving from 77% at M1 to 55% at M3 and 33% at M5. That is a meaningful structural signal. At the same time, the process detail shows that decision-heavy activities retain much higher human involvement than routine transactional activities. For example, candidate selection and offer negotiation remain highly human at M3, while applicant data management and other transactional activities fall much lower.

To me, that means the model is already telling an important story: AI does not simply remove humans. It changes where humans are required, what kind of work they retain, and how concentrated that work becomes.

2 Human Involvement Is Not the Same as Human Burden

A declining H% can easily be interpreted as a declining human workload. I am not sure that interpretation will always be true.

AI may remove high-volume, repetitive activity while leaving humans with work that is more ambiguous, emotionally demanding, cognitively intensive, exception-based, or consequential.

Human involvement can go down while human intensity goes up.

A recruiter may perform fewer manual activities but spend proportionally more time validating AI outputs, handling exceptions, explaining decisions, challenging questionable recommendations, and owning the final outcome. A manager may inherit fewer routine approvals but more escalations, judgment calls, and governance responsibilities.

That creates an important design question: are we measuring how much work remains human, or what it feels like and requires to successfully carry the human work that remains?

3 The Role-Level Dashboard Makes the Human Layer More Important

The role dashboard is powerful because it takes the process model and shows us who absorbs the change. The CHRO, HRBP Manager, Director of L&D, Compensation & Benefits Manager, and Director of HR Ops all show substantial movement across the maturity scale.

But an aggregate role percentage can hide the most important story. Two roles can have nearly identical H% values and still have completely different human consequences.

For example, an HRBP Manager and an HR Operations leader may both land in a similar M3 range, but the work they retain may differ dramatically:

Role	Likely Human Work That Remains	Potential Human Consequence
HRBP Manager	Judgment, coaching, influence, conflict, context, leader challenge	High relational and emotional intensity; psychological safety and authority matter
Director of HR Ops	Exceptions, controls, compliance, system/vendor oversight, escalation	High governance burden; risk concentration and capacity bottlenecks
Compensation & Benefits Manager	Fairness, market judgment, model interpretation, exceptions	High consequence of error; transparency and trust implications
Director of L&D	Capability strategy, learning design, adoption, behavior change	High transition burden; success depends on behavior, not only skill

Two roles can have the same H% and completely different Human Transition Risk.

4 I Think We Need to Separate Three Variables

The current H% is useful, but I wonder if the integrated methodology eventually needs to distinguish at least three concepts:

Human Involvement

How much human participation remains in the work?

Human Intensity

How demanding is the remaining work — cognitively, emotionally, relationally, and operationally?

Human Criticality

How consequential is it if the remaining human judgment fails?

These are not necessarily three separate indices. But keeping them conceptually distinct would help prevent a low H% from being interpreted as low human importance.

A role can be low in human activity volume and still be extremely high in human criticality.

5 The Maturity Scale Raises an Important Philosophical Question

The current maturity scale moves consistently toward lower human involvement as maturity increases. Visually, that can unintentionally communicate that “more mature” means “less human.”

I think we should challenge that assumption directly.

Should maturity mean less human involvement — or the right human involvement?

For transactional activity, greater automation may be the mature outcome. For judgment-heavy or high-consequence work, the mature outcome may be intentionally preserving human authority.

There may be a point below which reducing human involvement is no longer optimization. It may instead weaken judgment, accountability, trust, or control. That suggests a possible design concept:

Human Judgment Floor

For some responsibilities, there may be a minimum level of meaningful human participation below which the organization should not move, regardless of what technology is technically capable of doing.

The existing process patterns already seem to imply this. Decision activities retain much more human involvement than routine transactions. Making that principle explicit could make the maturity model stronger and easier to explain to executives.

6 Human Presence Is Not the Same as Human Agency

The phrase “human in the loop” can create a false sense of safety if the human is present but unable to meaningfully exercise judgment.

A person can technically remain in the workflow while still being:

- overloaded or rushed;
- undertrained on the AI system;
- unclear about decision authority;
- afraid to challenge the recommendation;
- unable to understand the basis of the output;
- conditioned to rubber-stamp machine recommendations;
- uncertain about when escalation is required.

Human presence does not guarantee meaningful human oversight.

This makes me think we should distinguish Human Presence from Human Agency.

Human Agency would ask whether the person has the authority, information, capacity, confidence, skill, and psychological safety necessary to exercise real judgment.

7 Psychological Safety May Become an AI Governance Variable

This is where the human layer becomes much more than a culture discussion.

If a recruiter believes an AI recommendation is wrong but does not feel safe challenging it, that is not simply an engagement problem. It is a decision-quality problem. If a manager fears that overriding the model will be interpreted as resistance, psychological safety becomes a governance issue.

Psychological safety may be a control mechanism in Human-AI work.

That would make the construct far more operational. It could help answer whether people are actually able to perform the human judgment the model intentionally leaves with them.

8 Workload Displacement May Really Be Complexity Displacement

Your “shock wave” idea made me think differently about workload.

AI may not merely move work. It may selectively remove the simple work and leave the difficult work behind. That can create an Exception Escalation Effect:

- routine work is automated;
- standard cases disappear from human queues;
- the remaining human cases are disproportionately ambiguous or difficult;
- exceptions move upward to managers or specialists;
- the average complexity of each human interaction increases.

A smaller volume of work can still create a larger management burden.

This is why “hours saved” may not equal “capacity created.” If the work that remains requires more judgment, more explanation, more governance, or more emotional labor, the organization may realize less capacity than the technical automation estimate suggests.

9 Accountability May Become More Concentrated as H% Declines

This may be one of the most important second-order effects.

As AI performs more execution, fewer humans may be involved in the workflow. But the final decision, override, exception handling, or accountability often remains human.

Activity can become distributed to technology while accountability becomes concentrated in fewer people.

That creates potential bottlenecks and risk concentration.

I wonder whether we should eventually examine an Accountability Concentration concept: as the human share declines, are we inadvertently placing more decision ownership, override responsibility, or reputational risk into a smaller number of human roles?

This connects directly to the principle we have already been using: AI can redistribute work. It cannot redistribute accountability.

10 The “No R/A Touches” Roles May Be Telling Us Something Important

The dashboard also shows roles with no current Responsible/Accountable touches in the underlying RACI data. I do not read that as evidence that those roles are unimportant.

It may instead expose a limitation of process-based visibility. Leadership roles often create value through activities that do not map neatly to formal R/A ownership:

- influencing;
- advising;
- integrating across functions;
- coaching;
- governing;
- sense-making;
- building trust;
- resolving conflict;
- interpreting organizational context.

The absence of an R/A touch may reveal invisible human work, not absent work.

As AI absorbs visible transactional activity, a larger share of human value may migrate toward work that is relational, interpretive, exception-based, and difficult to represent in traditional process architecture.

That may be exactly where our Human Transition layer is needed most.

11 Transition Distance May Need to Go Beyond Skills

The more I look at the role maturity view, the less convinced I am that the transition from current work to future work can be described as a skills gap alone.

A person's title may remain unchanged while the psychological nature of the job changes completely.

A recruiter, for example, may move from sourcing, screening, scheduling, documentation, and coordination toward validating, interpreting, challenging, persuading, managing exceptions, and owning decisions.

That suggests Transition Distance may need to capture:

- skill shift;
- judgment shift;
- identity shift;
- authority shift;
- relationship shift;
- accountability shift;
- work-pattern shift.

The real question may be: how far is the human being being asked to move from the way they previously created value?

12 A Possible Human Impact Pathway

Your shock-wave language may give us the connective spine for the integrated methodology.

Work Change → Role Impact → Human Impact → Manager Absorption → Team Effect → Accountability → Org Outcome

At each point in the pathway, we could ask a small set of consistent questions:

- What changed?
- Who absorbs it?
- Where did the work go?
- What new risk appeared?
- What new human capability is required?
- Who now owns the outcome?

I like this because it does not create a separate “HR layer” next to your architecture. It traces the human consequence through the architecture you already built.

13 A Potential Executive View

I can imagine the dashboard eventually telling two stories at once.

Story 1 — Structural AI Opportunity

Where can human involvement change? What does the HITL pattern look like by process and role?

Story 2 — Human Consequence

What happens to the person, manager, team, capacity, decision quality, and accountability because that involvement changed?

For an HRBP Manager, for example, the executive could see a maturity shift in H% and then a companion human-transition view such as:

Dimension	Illustrative Interpretation
Transition Distance	Reinvention
Human Criticality	High
Leadership Absorption	High
Workload/Complexity Displacement	Moderate to High
Psychological Safety Requirement	High
Accountability Concentration	High

The executive then sees not only where automation is possible, but what investment is required to absorb the change safely and effectively.

14 The Core Question I Think We Should Answer Together

Are we trying to optimize how little human involvement is required — or are we trying to design exactly where humans create the most value and make sure they can successfully carry the judgment, complexity, and accountability that remain?

If the objective is simply lower H%, the product risks being interpreted as an automation/workforce-efficiency model.

If the objective is intentional Human-AI allocation, the opportunity becomes much larger: an operating architecture that helps organizations determine what technology should do, where human judgment creates the most value, what conditions that judgment requires, and whether the organization can sustain the resulting design.

15 Questions I Would Like Us to Pressure-Test

- What exactly does a role-level H% represent, and what should we explicitly say it does not represent?
- Do we need to distinguish human involvement, human intensity, and human criticality?

- Should every maturity level move toward lower H%, or should some activities have an intentional human judgment floor?
- Can two roles with identical H% have different Human Transition Risk? If yes, how should we show that?
- What kinds of important human work are currently invisible because they do not appear as R/A touches?
- Should Consulted/Informed relationships or other leadership contributions be incorporated differently in future role analysis?
- How do we measure whether a human in the loop has meaningful agency rather than ceremonial presence?
- Where does workload move when transactional work disappears?
- Does automation create an exception escalation effect for managers?
- Does declining H% concentrate accountability into fewer human roles?
- How should Transition Distance account for changes in identity, authority, relationships, and accountability — not only skills?
- Which human-transition dimensions actually change an executive decision and which are interesting but unnecessary?
- What would make the combined model actionable enough for a CHRO, COO, CIO, or Chief AI Officer to use in a real workforce decision?

16 Organizational Role Taxonomy & Reporting Structure — Practical Sanity Check

There is a second piece I want to treat separately from the HITL maturity logic: the Organizational Role Taxonomy and the reportsTo relationships. You specifically flagged that several Managerial-tier reporting edges — particularly the HRBP Manager — were provisional after the self-consistency check against the RACI data. I agree that this deserves a practical operating-model review, not a cosmetic sign-off.

The RACI can tell us where responsibility and accountability appear in the process architecture. It cannot, by itself, prove that a reporting relationship is the most sensible one in practice.

What I want to sanity-check

- Does the proposed manager actually have the context and authority needed to lead the role?
- Does the reporting line reflect how decisions, escalation, coaching, performance management, and resource allocation work in practice?
- Is the relationship structurally sound across different organizational sizes, or is it only one plausible configuration?
- Does the proposed reportsTo edge create unnecessary layers, bottlenecks, or conflicts between functional expertise and business-facing accountability?
- Where should the taxonomy define a normative reporting line, and where should it explicitly allow common variants?

HRBP Manager — the edge I would examine most carefully

For the HRBP Manager, I would not validate the reporting line solely because it is internally consistent with the RACI. The role sits at the intersection of enterprise people strategy, business-unit context, employee relations, talent decisions, organizational design, and leader coaching. That makes the practical reporting relationship highly dependent on the operating model.

My working view is that the HRBP Manager usually needs a strong line into the senior HR/People leadership structure — for example a Head/Director/VP of HRBP, People & Culture, or CHRO structure depending on scale — while maintaining a very strong partnership with the business leader supported. A direct report into a business leader can create proximity, but it can also weaken HR independence, consistency, escalation, and enterprise-wide people governance if the model is not deliberately designed.

So the question is not simply “Who is closest to the work?” It is: “Who can provide the role with the right combination of functional stewardship, escalation authority, independence, coaching, and business access?”

Questions I would use to test every provisional reportsTo edge

- Who sets this role's priorities when business demands and enterprise people standards conflict?
- Who evaluates the role's performance and has enough visibility to do so credibly?
- Who coaches the role on the professional discipline, not only the business outcome?
- Who owns escalation when the role must challenge a senior leader?
- Who controls headcount, budget, systems, policy, and specialist resources the role depends on?
- Does the proposed manager have sufficient span, expertise, and organizational altitude to lead the role?
- Would this reporting line preserve independence in sensitive employee-relations, ethics, compensation, or workforce decisions?
- Does the role need a solid-line functional relationship and a dotted-line business relationship rather than a single simplistic edge?
- Would the answer change in a 500-person organization versus a 50,000-person enterprise?
- Is this relationship supported by the RACI because it is structurally correct, or because the underlying RACI itself reflects one organization's design assumptions?
- If AI changes the role's work, does the proposed manager become the logical absorber of the new exceptions, judgment, and accountability — or does the reporting structure need to change too?

Why this matters to the Human-AI architecture

I think reporting structure becomes even more important as AI changes work. If routine execution moves to technology while judgment, exceptions, escalation, and accountability become more concentrated, the manager above the role is not merely an org-chart box. That manager becomes part of the control architecture.

This means a reportsTo relationship that was workable in the current state may not automatically be the right relationship in the future state. Our role-transition analysis should therefore ask not only whether the role changes, but whether the managerial container around the role can absorb what the role becomes.

Possible design principle: AI-driven role redesign may require reporting-structure redesign when decision rights, exception burden, professional independence, or accountability migrate.

What I would like us to produce from the sanity check

- A confirmed primary reportsTo relationship for each prototype role where the evidence is strong.
- A clearly labeled “common variant” where multiple operating models are legitimate.
- A rationale for each relationship based on decision rights, escalation, professional stewardship, and operating reality — not only RACI proximity.
- A flag for relationships that should remain configurable rather than being presented as universal.
- A future-state check: whether the reporting line still holds after the role's HITL maturity and human-accountability profile change.

I would be comfortable telling you where I think the relationships hold in practice, where I think they are too context-dependent to hard-code, and where I think the future Human-AI operating model could actually require a different reporting relationship. That review feels substantive to me because it is exactly where structural consistency has to meet organizational reality.

17 Suggested Next Step

I do not think we should add more abstract constructs before testing the ones we already have.

I would suggest using the three prototype roles we have already discussed — Recruiter, HRBP Manager, and CHRO — and applying the same small set of questions to one representative responsibility in each role.

For each responsibility:

- What is the current human involvement?
- What changes under the target maturity state?
- What work disappears?
- What work moves?
- What becomes more complex?
- What human judgment remains?
- How critical is that judgment?
- Does the manager absorb new work?
- Does accountability become more concentrated?
- What must be true for the human judgment to succeed?
- What intervention would be required before implementation?

The goal of the prototype should not be to prove our methodology. It should be to break it, identify redundancy, expose blind spots, and determine what actually produces an executive-grade insight.

18 Where I Think Our Two Architectures Meet

The cleanest articulation I have right now is:

Your architecture helps determine where human judgment belongs.

My Human Transition layer helps determine what must be true for that judgment to work.

Together, the integrated methodology can determine how Human-AI work should be designed so that technology performs what it should, humans retain meaningful agency where it matters, and accountability remains intact across the system.

That feels to me like the intellectual center of what we are building.

19 Closing Thought

The most interesting thing about the maturity dashboard is not simply that human involvement declines. It is that the nature of human contribution changes.

As AI absorbs more routine execution, the human role may become smaller in volume but larger in consequence.

If we can show executives not only where work can be automated, but where human judgment becomes more concentrated, what pressure moves to managers, what invisible human work becomes more important, and what conditions are required for accountability to hold, I think we move beyond an automation model.

We begin to build a true Human-AI operating architecture.

Tim, these are the questions the current work is surfacing for me. I am not suggesting we adopt every concept above. I would rather pressure-test them with you and keep only what materially improves the architecture and the executive decision it is meant to support.

Source Context. These reflections are grounded in the current Human-AI Partnership Framework / PCF 7.0 HITL Maturity Assessment and the role-level HITL Design Dashboard reviewed in September 2026. The process assessment covers 37 Level 3 processes and 100 Level 4 activities across Develop and Manage Human Resources, with M3 identified as HITL Design Intent.

Working discussion document — [GiGi Garcia](#) to [Timothy King](#)

September 2026 · HITL-Driven Architecture · developed with the assistance of Claude (Anthropic)

[Download PDF](#) · [Back to WIP hub](#)