

EXTERNAL REVIEW · GIGI GARCIA · SEPTEMBER 2026

HITL Working Thoughts & Next Steps

My current read after the pressure test and our working discussion.

[Documentation](#)

[Dashboard](#)

[Work in Progress](#)

GiGi's own working notes, sent to Tim September 2026 after the working-relationship pressure test, transcribed here in the site's visual language so they can be read and referenced alongside the rest of the framework. The words and observations below are hers, not Tim's or Claude's — this page is a faithful transcription of her notes, not a response to them.

↓ Download PDF

Working note — not a finished position. This is where my thinking is landing right now. I want to keep pressure-testing it against real activities rather than turning every idea into another layer of the framework too quickly.

1 Where I Think We Are

The biggest shift for me is that the working relationship now feels like the most concrete place to start. It tells us, at the activity level, what the human and AI are actually doing together. That is much easier to understand and much easier to test than asking the pattern label to carry the whole story.

I would keep the original patterns in the backend because they may still matter analytically or technically, but I do not think the end user needs to see both. The working relationship is the thing that gives us something actionable: Automation, Agent-delegation, Augmentation, Oversight, Judgment, or Human-only.

From there, I think transition should be gated. Just because a technology can do more does not mean the work should automatically move. The way I am seeing it now, three things have to line up before a transition makes sense: the working relationship we actually want, the organization and people being ready to operate that way, and governance allowing that level of automation given the consequence of error.

2 What Feels More Settled

1. Working relationship should be the front door

I think this is the right pivot. It answers a practical question: how is this work supposed to happen between the human and the AI? It also gives us a cleaner way to talk about a process without making someone understand the

full technical model first.

What I still want to test is whether the six categories are distinct enough in real life. Augmentation versus Oversight is a good example. The definitions are different on paper, but I want to make sure a manager or employee could actually recognize which one they are operating in.

2. The patterns can stay in the data without being displayed

I do not want to throw them away. They may still be useful for the technical architecture, for consistency, or for future analysis. I just do not see a reason to ask an executive or practitioner to carry two taxonomies in their head at the same time.

3. Transition needs gates, not just a maturity curve

This is where the conversation became much clearer for me. A process should not move simply because AI capability exists. We need to know whether the proposed working relationship is appropriate, whether the human and organization are ready for it, and whether governance permits it. If any one of those fails, the transition should stop or remain at a lower level.

3 Where I Think I Can Add the Most Value

My strongest contribution is not trying to rebuild the technical architecture. It is helping define what has to be true on the human side before the architecture can actually work in practice.

Human Readiness / Agency

The pressure test keeps coming back to the same gap: a person can be named Accountable and still not be able to exercise that accountability. I want to make that visible.

The first version of the readiness gate does not need to be complicated. I would start by asking whether the human has:

- Authority — can they accept, reject, pause, or override the AI?
- Information — do they have enough context to make an independent judgment?
- Capability — do they have the knowledge and experience required for the remaining work?
- Capacity — do they realistically have the time and cognitive bandwidth to exercise judgment?
- Role clarity — is it obvious what the human owns and what the AI owns?
- Confidence — do they understand the technology well enough to challenge it?
- Psychological safety — can they question or override the AI without organizational penalty?
- Escalation / override — is there a clear path when the human cannot resolve the issue?

I would not rush to turn this into a fancy score. At first, I would use it as a structured conversation and mark each dimension as weak, adequate, or strong. The value is in seeing where readiness actually breaks down.

Human Criticality

I also think the consequence-of-error scale is a place where I can help make the framework more defensible. I would not replace Low / Moderate / High / Critical yet. I would first make sure we are asking the right questions underneath those labels.

- Who could be harmed and how?
- Is there legal or regulatory exposure?
- What is the financial or operational impact?
- Could trust or reputation be damaged?
- How wide is the reach — one person, a team, a workforce, or the enterprise?
- How reversible is the mistake once it happens?

Reversibility feels especially important. Some errors are immediately serious; others are slower, broader, and very difficult to unwind. Those are different risk shapes and the model should be able to explain the difference.

Human Intensity

This is the piece I do not want us to lose. AI may reduce volume without reducing the difficulty of the remaining human work. In HR, it may do the opposite: routine cases disappear first and the human is left with grievances, terminations, exceptions, ambiguity, and conflict.

I want us to ask not only, “How much human work remains?” but also, “What kind of human work remains?” I would initially look at cognitive complexity, emotional load, ambiguity, exception density, and decision stakes. That gets us away from assuming that hours saved automatically equal capacity gained.

Invisible human value

The RACI will never see everything that matters. Coaching, trust-building, influence, leadership, sense-making, and relationship repair may carry enormous value without appearing as Accountable or Responsible activity. I do not know yet whether this needs a formal axis, but I do think we need a way to flag when formal process data is likely to understate the real value of a role.

4 Where Tim Can Keep Driving the Architecture

I see Tim continuing to own the structure underneath the model: the data, activity classifications, rules, RACI logic, reporting relationships, and the technical mechanics of how the gates work together.

The areas I would want him to keep pushing are:

- Locking clear definitions for the six working relationships, including edge cases.
- Defining what evidence would move an activity from one working relationship to another.
- Keeping the original pattern taxonomy in the backend for analytical traceability.
- Creating a clean way to handle L4 exceptions when inherited L3 RACI does not reflect actual practice.
- Correcting provisional reporting relationships when the pressure test exposes better evidence.

- Building governance logic that can cap automation even when the technology could go further.
- Defining what every Governed Decision Point needs to include: owner, escalation, override, documentation, and review.
- Keeping version history so we can see why a classification changed and what evidence drove the change.

5 What I Think We Should Test Together

I do not think the next step is another large conceptual build. I think we should take a small set of very different activities and run them all the way through the same logic. If the method holds, great. If it breaks, that tells us exactly where to work.

I would use five activities because they stress different parts of the model:

- Select / reject candidate — high consequence, operational, AI can influence the option set before the final human decision.
- Manage employee grievance — high relational and emotional intensity, strong example of the Exception Escalation Effect.
- Define HR roles and accountability — executive judgment, broad organizational reach, difficult to unwind.
- Document employee performance — AI drafting opportunity with significant downstream legal risk.
- Establish succession plans — judgment-heavy, bias-sensitive, and broad organizational impact.

For each one, I would walk the same path: current activity → proposed working relationship → consequence / criticality → human readiness → governance limit → transition recommendation.

6 A Few Questions I Want to Keep Open

I do not think we need to answer all of these immediately, but I want to keep them visible because they may become important as we test the model.

- Are we classifying the activity itself, or the default working relationship for the normal case?
- Can the same activity have different working relationships depending on context or risk?
- Do we still need H% as a visible measure once the working relationship is clear, or does it create false precision?
- How do we account for AI influence when the human technically makes the final decision but AI shaped the options presented?
- When does reduced volume create capacity, and when does increased intensity consume that capacity?
- At what point has a role changed enough that we should call it a redesign rather than a transition?
- How do people develop judgment if AI removes the lower-risk work that historically taught them the job?
- What would we have to observe in the real world for us to admit one of our assumptions is wrong?

7 How I Would Start My Piece

I want my first contribution to be something we can actually use in the next working session, not another long narrative. I would start with three simple working pages.

Page 1 — Human Readiness Gate

The eight dimensions above, with a simple weak / adequate / strong read and a short note explaining the gap. The point is not to create a score for the sake of a score; it is to show why a technically possible transition may still fail.

Page 2 — Human Criticality Review

For each activity, walk human impact, legal/regulatory exposure, operational/financial impact, trust/reputation, reach, and reversibility. Then recommend Low / Moderate / High / Critical and explain why.

Page 3 — Three worked examples

One Recruiter example, one HRBP Manager example, and one CHRO example. Show the whole chain from activity through working relationship, readiness, governance, and final transition recommendation. If that reads clearly to us, we have something we can continue building.

8 What I Do Not Want This to Become

I do not want us to keep adding dimensions until the framework becomes impressive but difficult to use. Every new construct should earn its place by helping us make a decision, change a classification, set a control, or identify an intervention.

If it does not change what we would actually do, it probably does not need to be visible yet.

9 The Principle I Keep Coming Back To

The architecture is getting clearer to me as two sides of the same question. Tim is designing how the work should move. I am trying to make sure the human side can carry the work that remains.

The simplest version of that is:

Just because AI can perform the work does not mean the organization is ready — or should permit it — to do so.

That feels like the right place for us to keep building from.

Working notes — **GiGi Garcia**

Shared with **Tim King**, September 2026 · HITL-Driven Architecture · developed with the assistance of Claude (Anthropic)

[Download PDF](#) · **[Back to WIP hub](#)**