

THE HUMAN-AI PARTNERSHIP FRAMEWORK

Pattern Transition Narratives — Stage-First View

Role-Specific Pattern Transitions

What M1 → M3 → M5 actually means for the work — stage-first view, Recruiter

The treemap exhibits show how the composition of each role's remaining human work shifts as AI matures, organized by **maturity stage first, pattern second** — the same M1/M3/M5 column structure as the treemap itself. This view answers: at this specific stage, across every pattern the role touches, what does the work actually look like? Note that two different percentages appear below: the treemap's box size is each pattern's share of the total remaining workload at that stage (the boxes in each column sum to 100%), while the percentage next to each pattern in the stage detail is that pattern's own human-involvement level, independent of the others.


Recruiter

Accountable on 13 real L4 activities across Recruit/Source (7.2.2), Screen and Select (7.2.3), and Manage New Hire (7.2.4) — no Exception-pattern activities in this role's current RACI footprint.


As AI matures, Recruiter's remaining human work reshapes: Decision claims a growing share, Transaction shrinks toward the margins.

Composition of Recruiter's real activity mix still requiring human involvement, by HITL pattern — area = share, %


M1 — Undesigned
Current state



M3 — HITL Design Intent
Target state



M5 — Leading
Future state



Built from Recruiter's real PCF 7.2 RACI data (hitl_dashboard_final.html) — 13 L4 activities across its 3 Accountable processes (7.2.2 Recruit/Source, 7.2.3 Screen and select, 7.2.4 Manage new hire), weighted by activity count × each pattern's M1/M3/M5 human-involvement %. Recruiter has no Exception-pattern activities in current RACI data. Layout: squarified treemap (Bruls/Huizing/van Wijk), each column laid out independently from its own sorted values — box area is the only encoded quantity, position and aspect ratio are layout artifacts.

M1 — Undesigned

(Current state)

KNOWLEDGE · 85% human involvement

Determining recruitment channels, building and maintaining the candidate pool, choosing selection tools, and interviewing — at this stage the Recruiter owns essentially all of this judgment personally, with no structured AI support underneath it.

TRANSACTION · 60% human involvement

Recruiting events, vendor management, referral programs, channel administration, candidate testing, and hire processing — the Recruiter is personally running job fairs, updating vendor contracts, and processing the mechanics of a hire.

DECISION · 92% human involvement

Selecting and rejecting candidates, and negotiating the offer — the Recruiter's core identity. At this stage it's almost entirely a personal judgment call, with no structured AI input.

DOCUMENT · 72% human involvement

Drawing up and making the offer — largely manual drafting work at this stage, the kind of generative, templated writing AI absorbs quickly once introduced.

M3 — HITL Design Intent

(Target state)

KNOWLEDGE · 65% human involvement

AI is doing more of the pattern-matching underneath the work — surfacing likely-fit candidates, suggesting channel mix from historical fill-rate data, prepping interview guides — while the Recruiter still owns the synthesis: does this pool actually match what the hiring manager needs, is this candidate right beyond what a resume score shows.

TRANSACTION · 35% human involvement

Most of this work is now templated and system-driven — automated test administration, self-service vendor portals, workflow-triggered hire processing — with the Recruiter mainly handling exceptions the system flags.

DECISION · 80% human involvement

AI can flag background-check concerns or suggest a competitive offer range from comp data, but committing to hire someone and negotiating their terms stays fundamentally human. This is also where the role holds its two Governed Decision Points — Select/reject escalates to CHRO, Negotiate offer escalates to Director of TA — a real accountability boundary, not just a percentage.

DOCUMENT · 40% human involvement

Most of the language now auto-generates from approved comp data and role terms, with the Recruiter reviewing rather than drafting.

M5 — Leading

(Future state)

KNOWLEDGE · 42% human involvement

AI-assisted sourcing and structured interview scoring are the default; the Recruiter's value narrows to judgment calls AI can't make — reading real motivation, catching a mismatch the data missed, overriding a confident recommendation that's wrong.

TRANSACTION · 12% human involvement

This work is a rounding error at this stage: the Recruiter should rarely be touching it directly, reserved for the rare case the process genuinely breaks — the pattern that shrinks fastest and furthest of anything in the role.

DECISION · 65% human involvement

Even at the framework's most mature state, nearly two-thirds of the call to say yes or no to a specific person, and what it takes to close them, is still human — the highest floor of any pattern this role touches.

DOCUMENT · 18% human involvement

Nearly fully automated: the Recruiter's role becomes a final check that the specific terms match what was actually negotiated one pattern up, in Decision.

Source: [hitl_dashboard_final.html](#), PCF 7.2.2–7.2.4 (real RACI/pattern data) — matches [Recruiter_Job_Description.docx](#), [recruiter_treemap_mckinsey.html](#), and [pattern_transition_narratives.docx](#) (pattern-first companion view).